

Noise Removal?

Our story begins with image denoising ...




**Remove
Additive
Noise**



Denoising By Energy Minimization

$$f(\underline{x}) = \frac{1}{2} \|\underline{x} - \underline{y}\|_2^2 + G(\underline{x})$$



\underline{y} : Given measurements

\underline{x} : Unknown to be recovered

Denoising By Energy Minimization



$$f(\underline{x}) = \frac{1}{2} \|\underline{x} - \underline{y}\|_2^2$$

Relation to
measurements

$$+ G(\underline{x})$$

Prior or regularization

\underline{y} : Given measurements

\underline{x} : Unknown to be recovered

- ❑ This is in-fact a Bayesian point of view, adopting the Maximum-A-posteriori Probability (MAP) estimation.
- ❑ Clearly, the wisdom in such an approach is within the choice of the prior – **modeling the images** of interest.



Thomas Bayes
1702 - 1761

The Evolution of $G(\underline{x})$



During the past several decades we have made all sort of guesses about the prior $G(\underline{x})$ for images:

$$G(\underline{x}) = \lambda \|\underline{x}\|_2^2$$



Energy

$$G(\underline{x}) = \lambda \|\mathbf{L}\underline{x}\|_2^2$$



Smoothness

$$G(\underline{x}) = \lambda \|\mathbf{L}\underline{x}\|_{\mathbf{w}}^2$$



**Adapt+
Smooth**

$$G(\underline{x}) = \lambda \rho\{\mathbf{L}\underline{x}\}$$



**Robust
Statistics**

$$G(\underline{x}) = \lambda \|\nabla \underline{x}\|_1$$



**Total-
Variation**

$$G(\underline{x}) = \lambda \|\mathbf{W}\underline{x}\|_1$$



**Wavelet
Sparsity**

$$G(\underline{x}) = \lambda \|\underline{\alpha}\|_0^0$$

for $\underline{x} = \mathbf{D}\underline{\alpha}$





**Sparse &
Redundant**


- Hidden Markov Models,
- Compression algorithms as priors,
- ...



Sparse Modeling of Signals

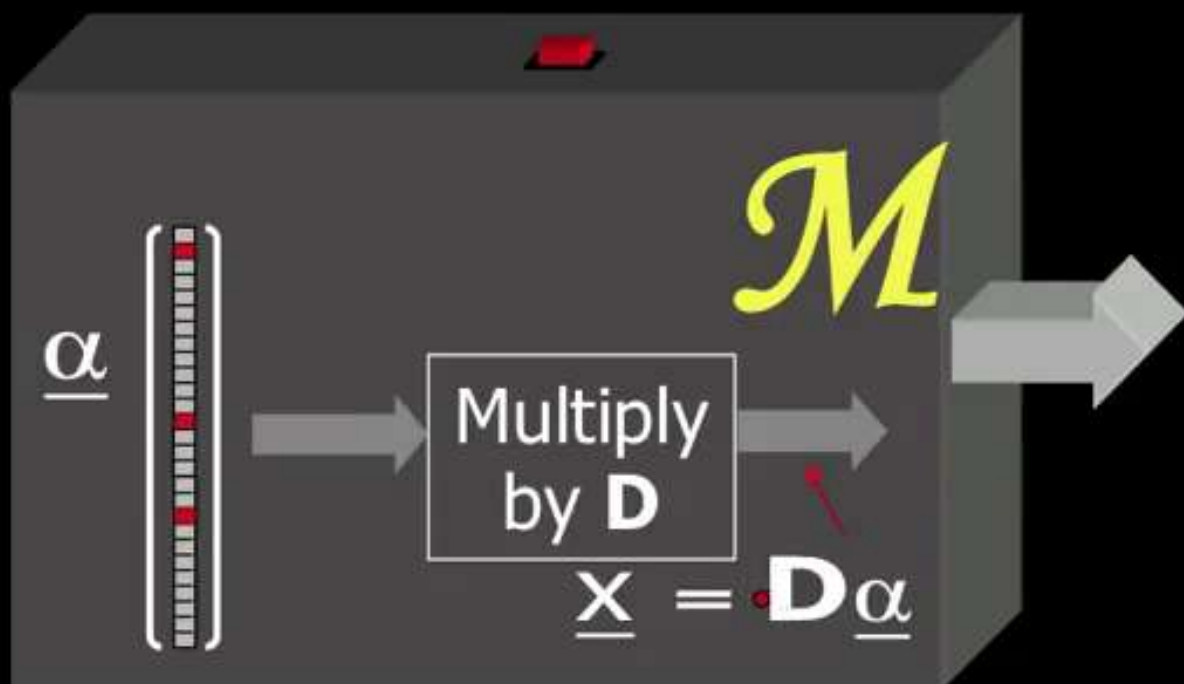


Sparse Modeling of Signals



- D (dictionary) is a set of prototype signals (atom).
- The vector α is generated with few (L) non-zeros.


Sparseland Signals are Special



- **Simple:** Every generated signal is built as a linear combination of **few atoms** from our **dictionary \mathbf{D}**
- **Rich:** A general model: the obtained signals are a **union of many low-dimensional spaces**.
- **Familiar:** We have been using this model in other context for a while now

JPEG

Sparse & Redundant Rep. Modeling




Our signal model $\underline{x} = \mathbf{D}\underline{\alpha}$ where $\underline{\alpha}$ is sparse

Sparse & Redundant Rep. Modeling



As $p \rightarrow 0$ we
get a count
of the non-zeros
in the vector

→ $\|\underline{\alpha}\|_0^0$



Our signal model $\underline{x} = \mathbf{D}\underline{\alpha}$ where $\|\underline{\alpha}\|_0^0 \leq L$

Back to Our MAP Energy Function



$$\frac{1}{2} \| \underline{x} - \underline{y} \|_2^2$$

Back to Our MAP Energy Function



$$\frac{1}{2} \|\mathbf{D}_{\underline{\alpha}} - \underline{y}\|_2^2$$

Back to Our MAP Energy Function



- L_0 "norm" is effectively counting the number of non-zeros in $\underline{\alpha}$.

$$\hat{\underline{\alpha}} = \arg \min_{\underline{\alpha}} \frac{1}{2} \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \text{ s.t. } \|\underline{\alpha}\|_0 \leq L$$

$$\hat{\underline{x}} = \mathbf{D}\hat{\underline{\alpha}}$$

$$\mathbf{D}\underline{\alpha} - \underline{y} =$$

Back to Our MAP Energy Function




- L_0 "norm" is effectively counting the number of non-zeros in $\underline{\alpha}$.

$$\hat{\underline{\alpha}} = \arg \min_{\underline{\alpha}} \frac{1}{2} \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \text{ s.t. } \|\underline{\alpha}\|_0 \leq L$$

$$\hat{\underline{x}} = \mathbf{D}\hat{\underline{\alpha}}$$

- The vector $\underline{\alpha}$ is the representation signal x .


$$\mathbf{D}\underline{\alpha} - \underline{y} =$$



- Few (L out of K) atoms can be combined to form the true signal, the noise cannot be fitted well. We obtain an effective projection of the noise onto a very low-dimensional space:
Denoising



Wait! There are Some Issues



- **Numerical Problems:** How should we solve or approximate the solution of the problem

$$\min_{\underline{\alpha}} \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \text{ s.t. } \|\underline{\alpha}\|_0^0 \leq L$$

$$\min_{\underline{\alpha}} \|\underline{\alpha}\|_0^0 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \leq \varepsilon^2$$

$$\min_{\underline{\alpha}} \lambda \|\underline{\alpha}\|_0^0 + \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2$$

- **Theoretical Problems:** Is there a unique sparse representation? If we are to approximate the solution somehow, how close will we get?
- **Practical Problems:** What dictionary \mathbf{D} should we use, such that all this leads to effective denoising? Will all this work in applications?

To Summarize So Far ...



Image denoising
(and many other
problems in image
processing) requires
a model for the
desired image

What do
we do?

We proposed a
model for signals/
images based on
sparse and
redundant
representation

Great!
No?

There are some issues:

1. Theoretical
2. How to approximate?
3. What about **D**?

Sparse Modeling: Some Theory and Implementation

Image and Video Processing: From Mars to Hollywood with a Stop at the Hospital

Guillermo Sapiro

Duke
UNIVERSITY


Li



Lets Start with the Noiseless Problem

Suppose we build a signal by the relation $\mathbf{D}\underline{\alpha} = \underline{x}$

We aim to find the signal's representation:



$$\hat{\underline{\alpha}} = \underset{\underline{\alpha}}{\text{ArgMin}} \|\underline{\alpha}\|_0^0 \quad \text{s.t.} \quad \underline{x} = \mathbf{D}\underline{\alpha}$$

Uniqueness

Why should we necessarily get $\hat{\underline{\alpha}} = \underline{\alpha}$?

It might happen that eventually $\|\hat{\underline{\alpha}}\|_0^0 < \|\underline{\alpha}\|_0^0$.

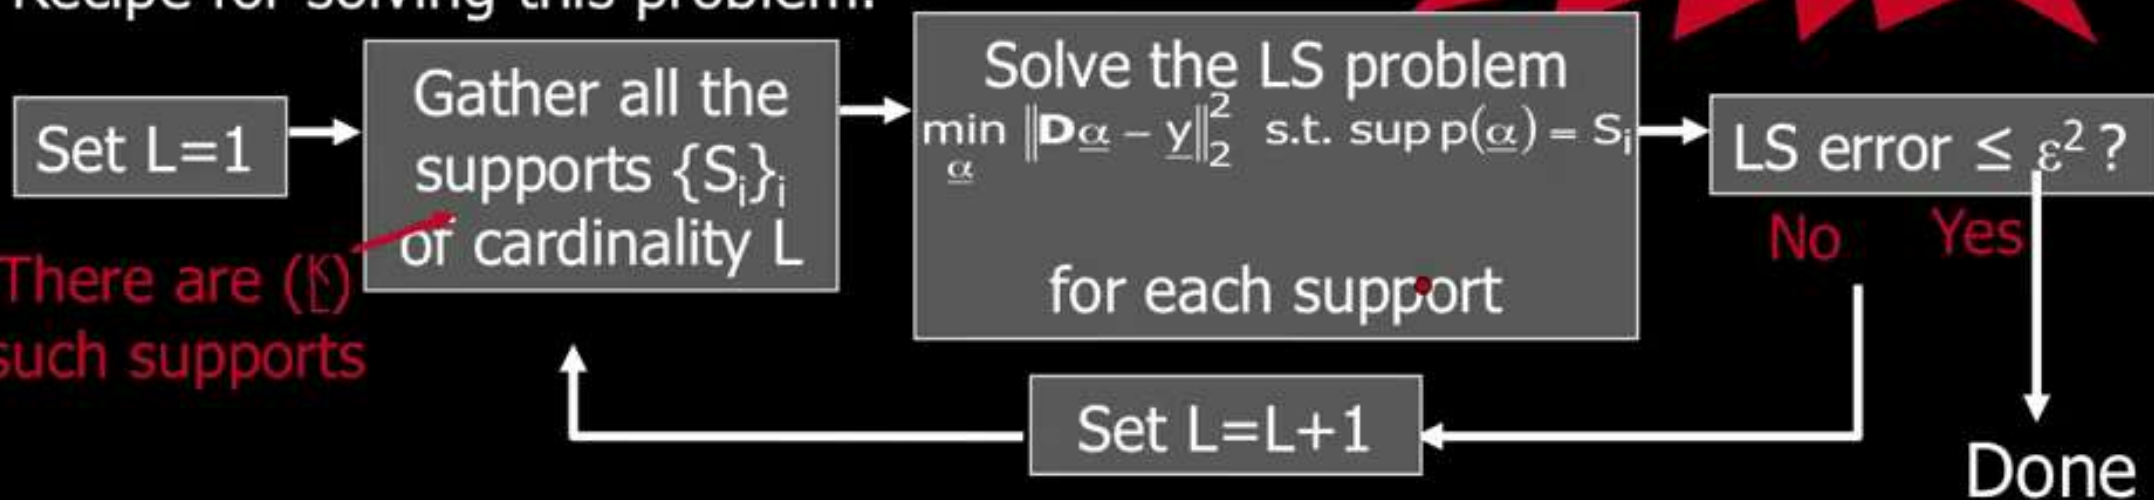


Our Goal

$$\min_{\underline{\alpha}} \|\underline{\alpha}\|_0^0 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \leq \varepsilon^2$$

This is a combinatorial problem, proven to be NP-Hard!

Recipe for solving this problem:



Assume: $K=1000$, $L=10$ (known!), 1 nano-sec per each LS

Our Goal

$$\min_{\underline{\alpha}} \|\underline{\alpha}\|_0^0 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \leq \varepsilon^2$$

This is a combinatorial problem, proven to be NP-Hard!

Recipe for solving this problem:

Set $L=1$

Gather all the supports $\{S_i\}_i$ of cardinality L

Solve the LS problem
 $\min_{\underline{\alpha}} \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \text{ s.t. } \text{supp}(\underline{\alpha}) = S_i$
for each support

LS error $\leq \varepsilon^2$?

No

Yes

Set $L=L+1$

Done

Assume: $K=1000$, $L=10$ (known!), 1 nano-sec per each LS

We shall need $\sim 8e+6$ years to solve this problem !!!!!

Lets Approximate



$$\min_{\underline{\alpha}} \|\underline{\alpha}\|_0^0 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \leq \varepsilon^2$$



Relaxation methods

Smooth the L_0 and use continuous optimization techniques

Greedy methods

Build the solution one non-zero element at a time

Relaxation – The Basis Pursuit (BP)



Instead of solving
 $\text{Min}_{\underline{\alpha}} \|\underline{\alpha}\|_0^0 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2 \leq \varepsilon$




Solve Instead
 $\text{Min}_{\underline{\alpha}} \|\underline{\alpha}\|_1 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2 \leq \varepsilon$



D, L



Relaxation – The Basis Pursuit (BP)



Instead of solving
$$\underset{\underline{\alpha}}{\text{Min}} \|\underline{\alpha}\|_0 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2 \leq \varepsilon$$



Solve Instead
$$\underset{\underline{\alpha}}{\text{Min}} \|\underline{\alpha}\|_1 \text{ s.t. } \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2 \leq \varepsilon$$

- ❑ This is known as the Basis-Pursuit (BP) [Chen, Donoho & Saunders ('95)].
- ❑ The newly defined problem is convex (quad. programming).
- ❑ Very efficient solvers can be deployed:
 - Interior point methods [Chen, Donoho, & Saunders ('95)] [Kim, Koh, Lustig, Boyd, & D. Gorinevsky ('07)].
 - Sequential shrinkage for union of ortho-bases [Bruce et.al. ('98)].
 - Iterative shrinkage [Figuerido & Nowak ('03)] [Daubechies, Defrise, & De-Mole ('04)] [E. ('05)] [E., Matalon, & Zibulevsky ('06)] [Beck & Teboulle ('09)] ...

Go Greedy: Matching Pursuit (MP)



- ❑ The MP is one of the greedy algorithms that finds one atom at a time [Mallat & Zhang ('93)].
- ❑ Step 1: find the one atom that **best matches** the signal.

$$\left[\begin{array}{c} \text{Matrix of atoms} \end{array} \right] \left[\begin{array}{c} \text{Selected atom} \end{array} \right] \approx \left[\begin{array}{c} \text{Signal} \end{array} \right]$$

$\|Dx - y\|^2$

Go Greedy: Matching Pursuit (MP)



- ❑ The MP is one of the greedy algorithms that finds one atom at a time [Mallat & Zhang ('93)].
- ❑ Step 1: find the one atom that **best matches** the signal.
- ❑ Next steps: given the previously found atoms, find the next **one** to **best fit** the residual.
- ❑ The algorithm stops when the error $\|\mathbf{D}\underline{\alpha} - \underline{y}\|_2$ is below the destination threshold.
- ❑ The Orthogonal MP (OMP) is an improved version that re-evaluates the coefficients by Least-Squares after each round.

Pursuit Algorithms



$$\min_{\underline{\alpha}} \|\underline{\alpha}\|_0^0 \quad \text{s.t.} \quad \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \leq \varepsilon^2$$

There are various algorithms designed for approximating the solution of this problem:

- ❑ Greedy Algorithms: Matching Pursuit, Orthogonal Matching Pursuit (OMP), Least-Squares-OMP, Weak Matching Pursuit, Block Matching Pursuit [1993-today].
- ❑ Relaxation Algorithms: Basis Pursuit (a.k.a. LASSO), Dnatzig Selector & numerical ways to handle them [1995-today].
- ❑ Hybrid Algorithms: StOMP, CoSaMP, Subspace Pursuit, Iterative Hard-Thresholding [2007-today].
- ❑ ...

Pursuit Algorithms



$$\min_{\underline{\alpha}} \|\underline{\alpha}\|_0 \quad \text{s.t.} \quad \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \leq \varepsilon^2$$

There are various algorithms designed for approximating the solution of

- Greedy Algorithms: Orthogonal Matching Pursuit (OMP), Least Squares Pursuit [2004-2006]
- Relaxation & numerical optimization
- Hybrid Algorithms: StOMP, CoSaMP, Subspace Pursuit, Iterative Hard-Thresholding [2007-today].
- ...

Why should they work?

To Summarize So Far ...



Image denoising
(and many other
problems in image
processing) requires
a model for the
desired image

What do
we do?

We proposed a
model for signals/
images based on
sparse and
redundant
representations

Problems?

The
Dictionary **D**
should be
found
somehow !!!

What's
next?

We have seen that there are
approximation methods to
find the sparsest solution,
and there are theoretical
results that guarantee their
success.

What Should **D** Be?



$$\hat{\underline{\alpha}} = \arg \min_{\underline{\alpha}} \|\underline{\alpha}\|_0^0 \quad \text{s.t.} \quad \frac{1}{2} \|\mathbf{D}\underline{\alpha} - \underline{y}\|_2^2 \leq \varepsilon^2 \quad \rightarrow \quad \hat{\underline{x}} = \mathbf{D}\hat{\underline{\alpha}}$$

Our Assumption: Good-behaved Images
have a sparse representation



D should be chosen such that it sparsifies the representations



One approach to choose **D** is from a
known set of transforms (Steerable
wavelet, Curvelet, Contourlets,
Bandlets, Shearlets ...)



Training

Learning from **Examples**

Measure of Quality for **D**


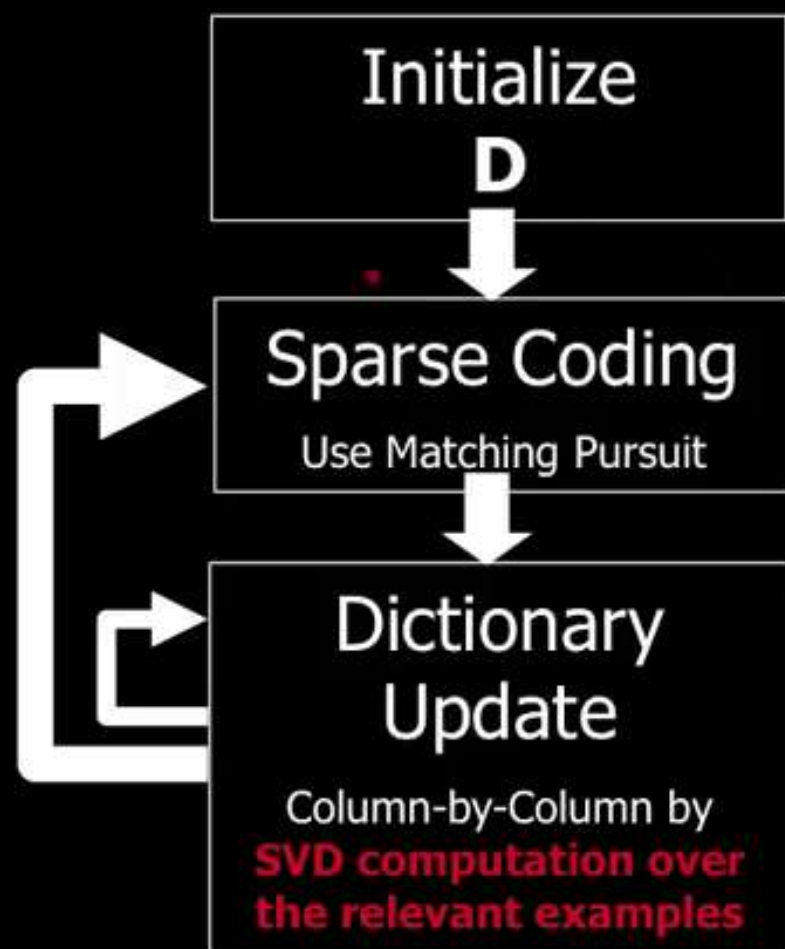


$$\underbrace{\begin{bmatrix} \text{X} & \dots \end{bmatrix}}_P \approx \begin{bmatrix} \text{D} \end{bmatrix} \begin{bmatrix} \text{A} & \dots \end{bmatrix}$$

$$\text{Min}_{\mathbf{D}, \mathbf{A}} \sum_{j=1}^P \|\mathbf{D} \underline{\alpha}_j - \underline{x}_j\|_2^2 \quad \text{s.t.} \quad \forall j, \|\underline{\alpha}_j\|_0^0 \leq L$$

on-line.

The K-SVD Algorithm – General



K-SVD: Sparse Coding Stage




$$\text{Min}_{\mathbf{A}} \sum_{j=1}^P \left\| \mathbf{D} \underline{\alpha}_j - \underline{x}_j \right\|_2^2 \quad \text{s.t.} \quad \forall j, \left\| \underline{\alpha}_j \right\|_p^p \leq L$$


D is known!
For the j^{th} item
we solve

$$\text{Min}_{\underline{\alpha}} \left\| \mathbf{D} \underline{\alpha} - \underline{x}_j \right\|_2^2 \quad \text{s.t.} \quad \left\| \underline{\alpha} \right\|_p^p \leq L$$


**Solved by
A Pursuit Algorithm**



K-SVD: Dictionary Update Stage




K-SVD: Dictionary Update Stage



We refer only to the examples that use the column \underline{d}_k

K-SVD: Dictionary Update Stage




We refer only to the examples that use the column \underline{d}_k



Fixing all \mathbf{A} and \mathbf{D} apart from the k^{th} column, and seek both \underline{d}_k and the k^{th} column in \mathbf{A} to better fit the **residual**!

K-SVD: Dictionary Update Stage



We refer only to the examples that use the column \underline{d}_k



Fixing all **A** and **D** apart from the k^{th} column, and seek both \underline{d}_k and the k^{th} column in **A** to better fit the **residual**!

K-SVD: Dictionary Update Stage



We should solve:

$$\text{Min}_{\underline{d}_k, \alpha_k} \left\| \alpha_k \underline{d}_k^T - \mathbf{E} \right\|_F^2$$

A diagram illustrating the minimization of the Frobenius norm of the residual matrix \mathbf{E} . The expression $\left\| \alpha_k \underline{d}_k^T - \mathbf{E} \right\|_F^2$ is shown. Arrows point from \underline{d}_k and α_k to a blue vertical bar, and from \mathbf{E} to a green grid.

We refer only to the examples that use the column \underline{d}_k



Fixing all \mathbf{A} and \mathbf{D} apart from the k^{th} column, and seek both \underline{d}_k and the k^{th} column in \mathbf{A} to better fit the **residual**!

K-SVD: Dictionary Update Stage



We should solve:

$$\min_{\underline{d}_k, \alpha_k} \left\| \alpha_k \underline{d}_k - \mathbf{E} \right\|_F^2$$

SVD

We refer only to the examples that use the column \underline{d}_k



Fixing all **A** and **D** apart from the k^{th} column, and seek both \underline{d}_k and the k^{th} column in **A** to better fit the **residual**!

To Summarize So Far ...



Image denoising
(and many other
problems in image
processing) requires
a model for the
desired image

What do
we do?

We proposed a
model for signals/
images based on
sparse and
redundant
representations

Problems?


Will it all
work in
applications?

What
next?

We have seen approximation
methods that find the
sparsest solution, and
theoretical results that
guarantee their success. We
also saw a way to learn **D**

$D \rightsquigarrow \text{subset } X$

From Local to Global Treatment



$$\hat{\underline{x}} = \underset{\underline{x}, \{\underline{\alpha}_{ij}\}_{ij}}{\text{ArgMin}} \quad \frac{1}{2} \|\underline{x} - \underline{y}\|_2^2 + \mu \sum_{ij} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}_{ij}\|_2^2$$

$$\text{s.t.} \quad \|\underline{\alpha}_{ij}\|_0^0 \leq L$$

From Local to Global Treatment



$$\hat{\underline{x}} = \underset{\underline{x}, \{\underline{\alpha}_{ij}\}_{ij}}{\text{ArgMin}} \quad \frac{1}{2} \|\underline{x} - \underline{y}\|_2^2 + \mu \sum_{ij} \left\| \mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}_{ij} \right\|_2^2$$

s.t. $\left\| \underline{\alpha}_{ij} \right\|_0^0 \leq L$

Extracts a patch
in the ij location

Our prior

What Data to Train On?



Option 1:

- ❑ Use a database of images



Option 2:

- ❑ Use the corrupted image itself !!



K-SVD Image Denoising



$$\hat{\underline{x}} = \underset{\underline{x}, \{\underline{\alpha}_{ij}\}_{ij}, \mathbf{D}}{\text{ArgMin}} \frac{1}{2} \|\underline{x} - \underline{y}\|_2^2 + \mu \sum_{ij} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}_{ij}\|_2^2 \quad \text{s.t.} \quad \|\underline{\alpha}_{ij}\|_0^0 \leq L$$

$\underline{x} = \underline{y}$ and \mathbf{D} known

Compute $\underline{\alpha}_{ij}$ per patch

$$\underline{\alpha}_{ij} = \underset{\underline{\alpha}}{\text{Min}} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}\|_2^2$$
$$\text{s.t.} \quad \|\underline{\alpha}\|_0^0 \leq L$$

using the matching pursuit

\underline{x} and $\underline{\alpha}_{ij}$ known

Compute \mathbf{D} to minimize

$$\underset{\mathbf{D}}{\text{Min}} \sum_{ij} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}_{ij}\|_2^2$$

using SVD, updating one column at a time

K-SVD Image Denoising



$$\hat{\underline{x}} = \underset{\underline{x}, \{\underline{\alpha}_{ij}\}_{ij}, \mathbf{D}}{\text{ArgMin}} \frac{1}{2} \|\underline{x} - \underline{y}\|_2^2 + \mu \sum_{ij} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}_{ij}\|_2^2 \quad \text{s.t.} \quad \|\underline{\alpha}_{ij}\|_0^0 \leq L$$

$\underline{x} = \underline{y}$ and \mathbf{D} known

\underline{x} and $\underline{\alpha}_{ij}$ known

Compute $\underline{\alpha}_{ij}$ per patch

$$\underline{\alpha}_{ij} = \underset{\underline{\alpha}}{\text{Min}} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}\|_2^2$$

$$\text{s.t.} \quad \|\underline{\alpha}\|_0^0 \leq L$$

Compute \mathbf{D} to minimize

$$\underset{\underline{\alpha}}{\text{Min}} \sum_{ij} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}\|_2^2$$

using the matching pursuit

using SVD, updating one column at a time

K-SVD

K-SVD Image Denoising



$$\hat{\underline{x}} = \underset{\underline{x}, \{\underline{\alpha}_{ij}\}_{ij}, \mathbf{D}}{\text{ArgMin}} \frac{1}{2} \|\underline{x} - \underline{y}\|_2^2 + \mu \sum_{ij} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}_{ij}\|_2^2 \text{ s.t. } \|\underline{\alpha}_{ij}\|_0^0 \leq L$$

$\underline{x} = \underline{y}$ and \mathbf{D} known

\underline{x} and $\underline{\alpha}_{ij}$ known

\mathbf{D} and $\underline{\alpha}_{ij}$ known

Compute $\underline{\alpha}_{ij}$ per patch

$$\underline{\alpha}_{ij} = \underset{\underline{\alpha}}{\text{Min}} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}\|_2^2$$

$$\text{s.t. } \|\underline{\alpha}\|_0^0 \leq L$$

Compute \mathbf{D} to minimize

$$\underset{\mathbf{D}}{\text{Min}} \sum_{ij} \|\mathbf{R}_{ij} \underline{x} - \mathbf{D} \underline{\alpha}_{ij}\|_2^2$$

Compute \underline{x} by

$$\underline{x} = \left[\mathbf{I} + \mu \sum_{ij} \mathbf{R}_{ij}^T \mathbf{R}_{ij} \right]^{-1} \left[\underline{y} + \mu \sum_{ij} \mathbf{R}_{ij}^T \mathbf{D} \underline{\alpha}_{ij} \right]$$

using the matching pursuit

using SVD, updating one column at a time

which is a simple averaging of shifted patches

$$\mathbf{R}_{ij} \underline{x} = \underline{x}_{ij} = \mathbf{D} \hat{\underline{\alpha}}$$


K-SVD




Image Denoising (Gray)



Source



Noisy image
 $\sigma = 20$



Initial dictionary (overcomplete
DCT) 64×256

Image Denoising (Gray)




Source



Result 30.829dB



Noisy image
 $\sigma = 20$




The obtained dictionary after
10 iterations

Denoising (Color)



- ❑ When turning to handle color images, the main difficulty is in defining the relation between the color layers – R, G, and B.
- ❑ The solution with the above algorithm is simple – consider 3D patches or 8-by-8 with the 3 color layers, and the dictionary will detect the proper relations.



Denoising (Color)



Original



Noisy (20.43dB)



Result (30.75dB)

Denoising (Color)



Original



Noisy (12.77dB)



Result (29.87dB)

Inpainting



Original



80% missing



Result

Inpainting



Original



80% missing



Result


Inpainting




Video Inpainting



Original




80% missing




Result


Video Inpainting



Original




80% missing



Result

Demosaicing



Side Note: Compressed-Sensing




- ❑ **Compressed Sensing** is leaning on the very same concepts, leading to alternative sampling/sensing theorems.
- ❑ Assume: the signal \underline{x} has been created by $\underline{x} = D\underline{\alpha}_0$ with very sparse $\underline{\alpha}_0$.
- ❑ Multiply this set of equations by the matrix Q which reduces the number of rows.

$$\underbrace{Q}_{N} \times \left\{ D \right\} = \left\{ x \right\}$$

Side Note: Compressed-Sensing



- ❑ **Compressed Sensing** is leaning on the very same concepts, leading to alternative sampling/sensing theorems.
- ❑ Assume: the signal \underline{x} has been created by $\underline{x} = D\underline{\alpha}_0$ with very sparse $\underline{\alpha}_0$.
- ❑ Multiply this set of equations by the matrix Q which reduces the number of rows.
- ❑ The new, smaller, system of equations is
$$\mathbf{QD}\underline{\alpha} = \mathbf{Q}\underline{x} \longrightarrow \tilde{\mathbf{D}}\underline{\alpha} = \tilde{\underline{x}}$$

- ❑ If $\underline{\alpha}_0$ was sparse enough, it will be the sparsest solution of the new system, thus, computing $D\underline{\alpha}_0$ recovers \underline{x} perfectly.
- ❑ Compressed sensing focuses on conditions for this to happen, guaranteeing such recovery.


Inverse Problems

$$\underline{\mathbf{y}} = \mathbf{U}\underline{\mathbf{f}} + \mathbf{w}$$
$$\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{Id})$$



\mathbf{f}

Inpainting



U masking

Deblurring



U convolution
 $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{Id})$

Zooming



U subsampling



Gaussian Mixture Models of Patches



$$\mathbf{y}_i = \mathbf{U}_i \mathbf{f}_i + \mathbf{w}_i$$


where

$$\mathbf{w}_i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}_d)$$

$$\frac{8 \times 8}{64}$$

- K Gaussian distributions or PCAs $\{\mathcal{N}(\mu_k, \Sigma_k)\}_{1 \leq k \leq K}$
- $\mathbf{f}_i \sim \mathcal{N}(\mu_k, \Sigma_k)$

$$K=10$$



Gaussian Mixture Models of Patches




$$\mathbf{y}_i = \mathbf{U}_i \mathbf{f}_i + \mathbf{w}_i$$

where

$$\mathbf{w}_i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}_d)$$



- Estimate $\{(\mu_k, \Sigma_k)\}_{1 \leq k \leq K}$ from $\{\mathbf{y}_i\}_{1 \leq i \leq I}$
- Identify the Gaussian k_i that generates $\mathbf{f}_i \forall i$
- Estimate $\tilde{\mathbf{f}}_i$ from $\mathcal{N}(\mu_{k_i}, \Sigma_{k_i}) \forall i$

Efficiently solved via MAP-EM





MAP-EM



Structured and Collaborative Sparsity



Sparse estimate



- Full degree of freedom in atom selection

$$\binom{K}{L} \sim 10^{14}$$

V.S.

Piecewise linear estimate



- Linear *collaborative* filtering in each basis.
- Nonlinear basis selection, degree of freedom $K \sim 10$.

Experiments: Inpainting



Zoom (original)



20% available 6.69 dB




PLE 30.07 dB



Experiments: Zooming



Motivation



Jogging



Running



Carrying



Jumping





Training

Class 1

Class 2

Class 3



Input Videos

Spatial Temporal Features

Sparse Modeling

Classification

l_1 Pooling



Classifier output



Results: YouTube Action Dataset

basketball	0.91	0.02	0.01	0.01	0.03	0.01	0	0.01	0	0	0
biking	0	0.97	0	0	0.03	0	0	0	0	0	0
diving	0	0	0.97	0	0.02	0	0	0	0	0.01	0
golf_swing	0.02	0.03	0.01	0.85	0	0.07	0	0.01	0	0.01	0
horse_riding	0	0.04	0.01	0	0.91	0.02	0	0	0	0.01	0.01
soccer_juggling	0	0	0.01	0	0.01	0.95	0	0	0.03	0	0
swing	0	0.08	0	0	0	0	0.92	0	0	0	0
tennis_swing	0	0	0	0	0	0.09	0	0.86	0.01	0.04	0
trampoline_jumping	0	0.03	0	0	0	0.03	0	0	0.92	0.02	0
volleyball_spiking	0.01	0	0.02	0.01	0.02	0	0	0	0	0.94	0
walking	0	0.02	0.02	0.01	0.03	0.01	0	0.01	0.01	0.01	0.89





